

Final Year Project Report

Automated Software to Understand Functional Relationship Between Dynamic Energy and Performance Events

Sukrat Kashyap

A thesis submitted in part fulfilment of the degree of

BSc. (Hons.) in Computer Science

Supervisor: Ravindranath Reddy Manumachu



UCD School of Computer Science

University College Dublin

April 6, 2018

Project Specification

General Information:

A energy model representing a relationship between dynamic energy consumption and performance events (PMCs) is constructed experimentally and the experimental dataset has the following format typically (k events, n records):

$$\begin{matrix} E_1, & x_{11}, & x_{12}, & x_{13} \dots x_{1k} \\ E_1, & x_{21}, & x_{22}, & x_{23} \dots x_{2k} \\ \dots & & & \\ E_n, & x_{n1}, & x_{n2}, & x_{n3} \dots x_{nk} \end{matrix}$$

where E_i is the experimentally obtained dynamic energy consumption of i-th record and x_{ij} are the experimentally obtained performance events (PMCs).

Given such an experimental dataset as an input, the goal is to determine/understand the functional relationship between the dynamic energy consumption and performance events (PMCs).

Two real-life datasets will be provided to the student.

Core:

The goal is to write automated software that will detect the following:

1. Existence of records where the dynamic energy consumption is the same (within an input tolerance) but all PMCs (with the exception of one) have same values. Then the relationship between energy and the one PMC is visualized to see the nature of the functional relationship.
2. Having accomplished step (1), understand the monotonicity of the relationship between dynamic energy consumption and performance events (PMCs).
3. Existence of records where the dynamic energy consumptions are different (within an input tolerance) but all PMCs have same values (within an input tolerance) suggesting the non-existence of a functional relationship.

The software must be written using any one mainstream language but preferably one of the following: C, C++, Python

The software must be well documented and tested.

Advanced:

Given an experimental energy model dataset as an input, the goal is to write software that performs intelligent but computationally feasible simulations where combinations of inputs are generated to study the existence/non-existence of a functional relationship between dynamic energy consumption and PMCs.

The software must be written using any one mainstream language but preferably one of the following: C, C++, Python.

The software must be well documented and tested.

Abstract

With the advent of technology, the demand of energy has increased dramatically. Increasing number of electrically driven equipment such as personal devices, hybrid vehicles and embedded systems have made the management of energy crucial. This projects focus is on understanding the energy consumption by High-Performance Computers. It has been known that CPU power consumption correlates to processor performance. And a number of models have been proposed to use performance events e.g. cache miss, DMA access, I/O interrupt, for the estimation of power consumption by the systems. This project takes a novel approach. Instead of estimating the energy consumption by the systems using performance events, a software is developed in order to better understand the functional relationship between the performance events and the energy consumption. A number of linear models were proposed due to the trickle-down effect of the performance events on power consumption [7]. The software also tries to analyze the linear dependence between the performance events and the energy consumption. This analysis helps validate the linear models that have already been proposed.

Acknowledgments

I would like to thank my supervisor Ravindranath Reddy Manumachu and mentor Alexey Lastovetsky for the assistance and guidance of this project. I would also like to thank my family and friends for supporting me in my work.

Table of Contents

1	Introduction	5
1.1	Motivation	5
1.2	Aims	6
1.3	Approach	7
1.4	Structure of report	8
2	Background Research	9
2.1	Energy consumption and Performance Events	9
2.2	Related Work	10
2.3	Existence of functional relation	10
2.4	Analysing functional relation	12
2.5	Clustering Methodologies	12
2.6	Software	15
2.7	Applications	16
3	Design and Optimisations	17
3.1	Design	17
3.2	Optimizations	19
3.3	Testing	21
4	Results and Evaluation	22
4.1	Results	22
4.2	Challenges	23
5	Conclusions and Future Work	24
5.1	Conclusion	24
5.2	Future work	24
A		26

Chapter 1: Introduction

1.1 Motivation

Modern day technology has developed under incredible speed in the recent decade and the computing power growth rate is truly phenomenal and lasting impact can be felt and benefit us in many ways. It is important to realise the worldwide effect on the environment by the increase in consumption of power by these technology advancements.

According to [10], the power consumption growth rates of PCs are about 7.5% per year. Data Centres and network play a much larger role as they both have power consumption rate of 12% each. This considerable growth is due to increasing data to be accessed, stored and processed. The constant expansion of energy consumption leads to increase in carbon emissions. CO_2 emissions from ICT (Information and communications technology) are increasing at a rate of 6% per year, at such rate by 2020 it will account to 12% of worldwide emissions [11].




Figure 1.1: Energy consumption distribution of data centers.

The major role of power consumption by a data warehouse is played by the servers and the cooling system which is used to cool down the physical parts of the server. The figure 1.1, shows that both the servers and the cooling systems account for 80% of the total consumption with both accounting for 40% each [11]. Making power and energy one of the key challenges in system and software designs. In order to improve processor's performance while limiting power consumption, designers have increasingly utilized dynamic hardware adaptations. But to guide to these hardware adaptations it is necessary to measure the power of systems accurately. Such adaptations can reduce power consumption and chip temperature by reducing the amount of available performance. Temperature sensors are slow in response due to the thermal inertia of the microprocessor. Relying on them would give slow response and detection of temperature changes. Many demonstrations have been done to show that performance counter/performance events are effective proxies of power measurement [3].

The study of the relations between different performance counters and energy consumption has

become crucial. The existence and analysis of the functional relation between them will enable minimization of power consumption leading to less heat generation from the CPU. This decrease in temperature will further reduce the usage of the cooling system. Which all, in turn, leads to cost minimization and a step towards the eco-friendly environment.

Thus, the discovery of the functional relationship between the power consumption and the performance events (PMCs) of the system, will enable the ability to predict the energy consumption for certain computations.

1.2 Aims

The project aims to create a piece of software that will help to understand and find relations between energy consumption and performance counters. The software must aim to be extensible, easy to use and performant. It should be able to perform analysis like regression and clustering on the data. These actions will help in showcasing the relationship of various parameters and the output. The parameters are the performance events and output is the energy consumption by the systems. Following are the objectives that are attempted to be accomplished.

1.2.1 Existence of functional relation

Datasets are mostly in pair of multiple inputs and a single output. And the possibility of the existence of a relation is what has to be found. Is it possible to define the data in a form of function? A question like this is one of the aims. Defining the data in a form of function can be thought of as explaining the dataset in a form of a formula which can combine various parts of the input parameters and match the output. It is quite impossible to tell that whether a function definitely exists. This is because the data is always a subset of the population and there are a number of records which are either not in the dataset or it is not known. But, regardless it can definitely explain the non-existence of a functional relation by finding a pair of input parameters and output that violate the definition of a function.

The pair which violates the definition of a function must be looked at as there is a high probability of that data set record being corrupt or if not it can actually act as a contradiction and explain the non-existence. If this outlier is, in fact, corrupt, this process can be used for data cleanups as well.

1.2.2 Analysis of functional relation

Once, it is not possible to prove functional non-existence in the dataset anymore. The dataset must be analysed in order to better understand the functional relation of the parameters and the output. The analysis of the relation explains the contribution of a parameter to the result. It gives a better understanding of the correlation between two quantities. Does high page faults correlates to high energy consumption? Questions like this are what the analysis tries to answer. This can also be thought of an explanation of the resultant parameter.

In other words, understanding the correlation is the other aim. Once it is known that there is a high correlation. Further investigation can be done in order to find the form of the relation.

1.3 Approach

For each aim, two approaches have been employed. Both the approaches try to achieve the same conclusions but with different methodology.

The experimental data sets provided has the following format:

$$\begin{aligned} E_1, x_{11}, x_{12}, x_{13} \dots x_{1k} \\ E_1, x_{21}, x_{22}, x_{23} \dots x_{2k} \\ \dots \\ E_n, x_{n1}, x_{n2}, x_{n3} \dots x_{nk} \end{aligned}$$

where k is the number of parameters (performance events), n is the number of records in the dataset. E_i is the experimentally obtained dynamic energy consumption and x_{ij} are the experimentally obtained performance events (PMCs).

1.3.1 Existence of functional relation

The main objective here is to find the non-existence of functional relationship. In other words, it means proving that the dataset cannot be explained in terms of a function/formula.

First approach here is to find two performance events tuples $(x_{i1}, x_{i2}, x_{i3} \dots x_{ik})$ and $(x_{j1}, x_{j2}, x_{j3} \dots x_{jk})$ that are equal within some tolerance, but their corresponding E_i and E_j dynamic energy consumption are different. Existence of such a tuple in database will lead us to prove the non-existence of a functional relationship.

It is, in fact, easier to find two equal records by ordering the dataset by their parameters and going through the entire dataset once to find equal records and comparing their dynamic energy consumption. The order of complexity of such is $O(N \log(N))$ which is the complexity of sorting as that is the only heavy duty task involved. But it gets complicated when tolerance comes into play. Likewise, the task at hand is to find similar records rather than equal records. The project employs 2 methods to measure the similarity between the two data records.

First approach:

The data records are imagined as data points in $k - space$ as we have k number of parameters. Then the whole space is divided into small $k - dimension$ cubes with dimensions $(t_1, t_2, \dots t_k)$ where t_i is the tolerance of each parameter. The data points are then put in their respective cubes. When two points lie in the same cube that means their parameters are similar and so must be there output. If they do not have similar output within some tolerance then they violate the functional relation.

Second approach:

The data points in $k - space$ are said to be similar if the Euclidean distance between them is less than the t provided, where t is the total/max tolerance. In this method, clustering is done using the distance between the parameter vector of the data points.

1.3.2 Analysis of functional relation

Many different types of relations can exist between two variables. But, since the correlation between the variables is known to be linear. Linear regression is done between the dynamic energy consumption with the performance events. Many researchers have been successful in estimating using linear models to estimate energy consumption using performance events. [9] The preference given to linear models is due to the trickle-down effect of the performance events. [4] Hence, Linear regression is used to analyse the relationship.

Linear regression is done with 2 variables but there are k variables related to energy consumption making it hard to understand the nature. Multiple linear regression tries to find the best fit to the data but in this process, the best fit hinders the tolerance factor in the dataset and also undermines the actual relation. As multiple regression is also interested in predicting the output variable. Hence to analyse one of the variables, the data point must be grouped using their other $k - 1$ variables. Each grouped cluster is then visited and linear regression is performed. Two approaches are employed to cluster the dataset. These are similar to clustering techniques used in finding the existence of the relation.

First approach:

Data points are clustered by putting the data points in the respective $k - 1$ dimension cube. Forming a cluster with almost similar $k - 1$ variables. Each data point is put in one of the cubes. These data points have similar $k - 1$ performance event vector but varying one of the parameters and output variable. Regression line calculated corresponds to the effect of the parameter variable to the output variable in a particular situation. The situation is the similar $k - 1$ parameter values.

Second approach:

Euclidean distance is used to isolate similar $k - 1$ variables. A $k - 1$ dimension sphere can be imagined for simplicity with its radius being the total/max tolerance. Here, each data point has its own small sphere. The regression line calculated for each data point act as the tangent to the data point in the direction of its neighbours.

1.4 Structure of report

In this report, the motivation behind the project and brief description of the approaches that will be undertaken are seen already.

This Introductory chapter is followed by Background research, design aspects of the software, implementation of the software, testing and evaluation of the tool followed by the conclusion and the future works.

Background research explains about performance events and how do they influence energy consumption by the system. It also explains the approaches in detail mentioned above and analyses their complexity and how they can be optimised.

Following Background research, design and implementation of the software are deep dig into as we want an extensible, performant tool. Testing and evaluation of the software are also explained. How the software is tested and its evaluation of its performance. The report ends with conclusions and future works that shows how the project can be extended and applied to various other domains.

Chapter 2: Background Research

Many designers are increasingly utilizing dynamic hardware adaptations to improve performance while limiting the power consumption. Whereas Some are using software to decrease power usage for e.g. putting the system in sleep mode when it's in the idle state. The main goal remains the same, which is to extract maximum performance while minimizing the temperature and power. The study and examination of the events affecting the consumption will help us to predict and minimize the consumption of energy with very high accuracy.

2.1 Energy consumption and Performance Events

First, let's look at energy consumption. Energy consumption is the power (usually in watts) consumed by a system. This system could be the processor/CPU, memory, disk, I/O (Input/Output) system, chipset or the whole computer system itself. It has been known that power consumption of any system has a high correlation to its usage. Since the correlation is high for its usage. It means usage is a good way to understand the power consumption. Performance events are one of the better ways to measure the usage of a system as they are known much more faster than the temperature sensors. As temperature sensors are slow in response due to the thermal inertia of the microprocessor.

Now let's look at what are performance events, performance events are any events that affect the consumption of energy in some way. Selection of performance events is quite challenging. A simple example would be the effect of cache misses in the processor. For a typical processor, the highest level of cache would be L3 or L2 depending on the type of processor. Now for some transaction which could not be found in the highest level of cache (cache miss) would cause a cache block size access to the main memory. Thus, the number of main memory access would be directly proportional to the cache misses. Since this memory access is off-chip, power is consumed in the memory controller and DRAM. Even though, the relation is not simple as it seems but a strong casual linear relationship between the cache miss and the main memory power consumption [4].

Another performance event can be instructions executed. The more instruction being executed, will turn on and use more units of the system. Hence, power is consumed as opposed to when the processor is in its idle state [7].

Cache miss, TLB misses are also good performance events as they seem to have a strong relationship between the power consumption as processor needs to handle memory page walks. Same can be said for Page faults where a program is not able to find mapped address in physical memory as it has not been loaded yet. This causes a trap which can result in a number of situations, one of them which is to get the data from disk. It involves a long walk to the disk. This walk to the disk and raising of exceptions would consume more energy by the disk as well as the CPU. Another thing to note here most of the relations that we saw above are directly proportional to each other. Increase in the variable like cache miss, the number of cycles in CPU etc gives rise to energy consumption by the system. This makes a linear model a good point to start when analysing the relationship between the performance events variables and energy consumption.

2.2 Related Work

In this section, some of the prior research are discussed. Hardware performance counter's links to processor power consumption were first demonstrated by Bellosa. In [3], Bellosa demonstrated the high correlation between performance counters such as memory references, L2 cache, floating point operations to processor power consumption.

Gilbert in [7], predicted the power consumption for Intel XScale processors using performance monitoring unit events. Since power consumption is greatly dependent on executing workload, power estimation was done using HPCs (Hardware performance counters) such as Instruction cache miss, TLB misses etc. The linear parameterisation of power consumption based on performance events was done based on performance events.

In [13], proposed a full system power model for CPU-intensive and memory intensive applications with active cycles, instruction retired and LLC missed as performance events. A full power model for I/O intensive applications was also proposed considering the system level utilization as performance events. Many machine learning based algorithms like logistic regression, elastic net and k-nearest neighbours were applied to the real-world application.

Bircher approached in a distinct way by using events local to the processor and eliminating the need for sensors spread across various parts of the systems [5]. Linear regression modelling was done in order to predict the power consumption at runtime. Multiple linear and polynomial regression was done only when accuracy was not obtained.

The high correlation between performance events and power consumption was demonstrated by all of them. But all of them tried to predict the power consumption using system events. They use various methodologies in predicting so. In [13], a number of machine learning algorithms were used. But in this project, an attempt is made to understand the monotonic relation between the events and power. The predictive models would work on a particular specification of the system. But understanding the relation will help define a model on the architecture level. This will also help in verifying the models that already exist.

2.3 Existence of functional relation

Definition: Given a dataset of pairs (x_i, y_i) where $i \in [1, n]$ of two variables x and y , and the range X of x , y is a function of x iff for each $x_0 \in X$, there is exactly one value of y , say y_0 , such that (x_0, y_0) is in dataset. [14]

Above is the definition of functional relation. In our case the $x_i = (p_1, p_2, \dots, p_k)$ and $y_i = E_i$ where p are the performance events and k is the number of events and E is the energy consumption.

In other words, functional relationship is an one to one mapping between our input variables p_1, p_2, \dots, p_k and output E_i . This reasoning can be explained quite intuitively as assuming functional relation we can formulate a $f(p_1, p_2, \dots, p_k) = E_i$ where f is function. Now if one (p_1, p_2, \dots, p_k) can give more than one output E_i then that f function is either not correct or f does not exist. A question that immediately arises is what if two different (p_1, p_2, \dots, p_k) gives same output E_i . The answer is it is possible and it does not violate the functional relation definition as there is still 1 to 1 mapping from input to output. The only difference is the functional relation is not surjective anymore which means that one cannot figure out the input values from output values (i.e. the other way around). But we are only interested in predicting the output

rather than inputs from the output variables.

The Proof of the existence of functional relation is given below:

But first, let's look at our dataset that will be provided. We know that the data will be in the following format:

Let k be the number of parameters for the energy and n be the number of records in the dataset

$$E_1, x_{11}, x_{12}, \dots x_{1k}$$

$$E_2, x_{21}, x_{22}, \dots x_{2k}$$

...

$$E_n, x_{n1}, x_{n2}, \dots x_{nk}$$

where E_n is the dynamic energy for the n th tuple and x_{nk} corresponds to the k th performance event for n th record.

We will use mathematical definition of functional relationship given above to prove the approaches:

To Prove: We need to prove that finding at least 2 equal performance events with different dynamic energies ensures that there exists no functional relationship in the dataset.

Proof. Let us assume that there exists a functional relation such that:

$$f(x_{n1}, x_{n2}, \dots x_{nk}) = E_n$$

where f is the functional relation for the dataset.

If we have $f(x_{i1}, x_{i2}, \dots x_{ik}) = E_i$ and $f(x_{j1}, x_{j2}, \dots x_{jk}) = E_j$ where $E_i \neq E_j$ and $(x_{i1}, x_{i2}, \dots x_{ik}) = (x_{j1}, x_{j2}, \dots x_{jk})$

If such i and j exists. Then, we can conclude that the f is not a function by using the definition of a function as this assumed function has two images.

Which contradicts the hypothesis stated above. Hence by proof of contradiction, we could say that f is not a function of the dataset. \square

Now the task is to find similar $x_{11}, x_{12}, \dots x_{1k}$ input variables. If it was equality, it could be performed trivially by sorting the data records on the basis of $x_{11}, x_{12}, \dots x_{1k}$. Followed by going through the data records in that order to find equal records. As equal records will be next to each other when sorted.

But the dataset accumulated is collected from experimental setup. Experimental setups data have some error/tolerance associated with it. As the equipment or software cannot accurately measure and have some error associated with it. e.g. Energy consumption is measured by a power meter which does have a confidence interval. The tolerance associated with the dataset no longer allows performing equality on input variables. Hence, a method is needed to measure the equality of the tolerance.

Now after knowing the tolerances the data records in the dataset now looks more of the form:

$$E_n \pm e_E, x_{n1} \pm e_1, x_{n2} \pm e_2, \dots x_{nk} \pm e_n$$

where e is now the error associated with the variable.

Hence dataset must be clustered with the tolerances associated with them. Data points having similar k parameters will fall in the same cluster. The output variable $E_n \pm e_E$ must be close to each other within some tolerance for all the points in the cluster. If they are not, it violates the functional relation and the data points are picked and must be given to the user to either discard if corrupt. This removal of corrupt data records is known as Data cleanups. If they are not

corrupt then it proves the non-functional relation in the dataset. Analysis of functional relation can be done when no records are found that contradicts the functional relation definition.

2.4 Analysing functional relation

Now this section corresponds to analysing the functional relation. A functional relation can of many forms, there can be logarithmic functional relation, exponential, linear, polynomial etc. According to [4], shows that functional relation between performance event and dynamic energy is of the linear form. This is due to the trickle-down effect of the performance events. Hence, Linear model is used to understand the relationship between the events and the consumption.

This project is more interested in finding whether a strong relation of monotonicity exists in the dataset or not. And linear regression is best suited for this kind of task. Linear regression is an approach determine how strongly two variables correlate with each other. Correlation does not necessarily mean causation. Looking at the data points and linear regression values such as Pearson coefficient, R^2 can determine how strong the correlation is between 2 values but it cannot explain the cause of it.

The dataset contains k parameter variables, to analyse the relation between the parameter variables and the output variable. One parameter is chosen at a time. To see the change in the parameter variable chosen and the output variable, isolation of the other variables is required. Isolation is done by grouping the data by $k - 1$ parameter variable forming a number of clusters with similar $k - 1$ parameters. These clusters are then visited and linear regression is performed on the k^{th} variable and the output variable.

From the steps above, it can be observed the parameter to be analysed is isolated by grouping the dataset with the variables that are not being analysed and are equivalent. Since clustering of the dataset is performed again but with different vector dimension of $k - 1$, the clustering algorithms are needed by both the objectives.

2.5 Clustering Methodologies

This section introduces two clustering algorithm that was used to cluster data points which are similar or close to each other needed by both of the above objectives. There are many clustering methodologies out there [12]. No clustering algorithm can universally solve all the problems. An algorithm which favours certain observations, assumptions and favours some type of biases are designed and used.

Clustering involves grouping similar data by their attributes. There was a number of clustering algorithms like Hierarchical clustering, K means clustering, Graph-based clustering e.g. Chameleon. But we chose Grid-based clustering and Distance-based clustering.

Most of the clustering algorithms like K means clustering requires the number of clusters to be formed. In the dataset, the number of clusters needs to be formed is not known. In Hierarchical clustering, the hierarchy of dataset is used but the relation is functional and not hierarchical. Graph-based clustering usually requires edges, and edges could be formed but they would increase the space complexity exponential. As more storage will be required to store the edges of the

dataset.

The two types of clustering algorithms chosen use some facts about the dataset. Datasets with tolerances for clustering would be best suited to use these algorithms.

2.5.1 Grid based clustering

Since tolerance must be used as a measure of clustering data points. The definition of equality for a variable changes from $x_{ia} = x_{ja}$ to $|x_{ia} - x_{ja}| \leq e_a$. Now, if simple sort and search for similar variable would be employed, it would not work correctly as equivalent records would not be next to each other. They might non-equivalent records in between.

Example data to illustrate the same:

E	x_1	x_2
3.5	4.5	6.5
4.1	4.6	10.6
0.2	4.7	6.5
1.6	4.7	7.6

The above three records are sorted by their parameters x_1, x_2 . We can see that if the $e = 0.5$ is the absolute error. Then record number 1 and 3 are similar to each other but do not lie next to each other. This increases the complexity of finding similar records from $O(N \log(N))$ to $O(N^2)$ as we do not know where the similar records will lie, so N into N search must be performed which is quite inefficient.

To make it efficient, instead of finding similar records in the dataset. The whole $k - dimensional$ space is divided into small $k - cubes$ whose dimensions are $(e_1 * e_2 * \dots * e_k)$. Each $k - cube$ has its own integer coordinates in space. The data records are grouped together with respect to their respective $k - cube$ coordinates. Since the coordinates are integers (equality can be performed). They can be grouped in $O(N \log(N))$ time. The calculation of the coordinates in which the data point belongs to can be calculated in $O(1)$ time by the following.

$$coordinate = \left(\left\lfloor \frac{x_{n1}}{e_1} \right\rfloor, \left\lfloor \frac{x_{n2}}{e_2} \right\rfloor, \dots, \left\lfloor \frac{x_{nk}}{e_k} \right\rfloor \right)$$

Every data point will have a corresponding coordinate they belong to. Data points are then grouped together by their coordinates. Each and every grouped coordinate is a $k - dimensional$ cube.

Algorithm 1 Grid based clustering

```

1: procedure GRID( $a, b$ )
2:   Input dataset: a list of data points,  $e$ : errors for each variable
3:   Result clusters created with each cluster having its unique index
4:   list  $\leftarrow$  List.empty
5:   for each point in dataset do                                      $\triangleright N$  iterations
6:     coordinate  $\leftarrow$  ( $\lfloor point.x_1/e_1 \rfloor, \dots, \lfloor point.x_k/e_k \rfloor$ )
7:     list.add( $[coordinate, point]$ )
8:   sort list by the coordinate value                                    $\triangleright N \log(N)$  for sorting
9:   result  $\leftarrow$  group list by the coordinate value                  $\triangleright N$  iterations
10:  return result

```

It can be seen since space is divided into cubes of a certain dimension. The cubes are disjoint and they do not overlap. It can be imagined as a building of boxes stack on top and side of each other. There will be points which lie on the edges of the n -cube. Since the cubes do not overlap. They will be close to some of the points in the cubes next to them. But in the algorithm, one data point is assigned to only one cluster. This problem is solved by the next approach which is the distance based clustering. This problem comes with a performance advantage.

The algorithm complexity can be measured.

- Time complexity $O(N \log(N))$: This is because of sorting which is the slowest operation in the algorithm, but nevertheless optimizations can be done by using a Dictionary or Binary search tree to group data with equal coordinates. Using Dictionary will give the complexity of $O(N)$.
- Space complexity $O(N)$: For each data point only the corresponding *coordinate* is stored.

2.5.2 Distance based clustering

In this method, data points are thought of as vectors in Euclidean space of n -dimension where n is the number of parameters considered when clustering. The equality of vectors according to which clustering should take place is done by using the Euclidean distance between two vectors. Two vectors are said to be equivalent if the distance between the two vectors is less than the tolerance specified.

Let u and v be two vectors in space with dimension n , then the euclidean distance is define as:

$$\text{dist}(u, v) = \sqrt{(u[1] - v[1])^2 + (u[2] - v[2])^2 + \dots (u[n] - v[n])^2}$$

Now, two vectors u and v are said to be equal or close enough when:

$$\text{dist}(u, v) \leq \text{tol}$$

where tol is the maximum distance between two points to call them neighbours.

This tolerance can be thought of as the total tolerance of the between two data points. Because the difference between two points in all of their different dimensions are squared and “added” together and then square rooted. It signifies the total tolerance that is allowed between two data points.

This clustering is creating spheres of k -dimension for each data point. Here, overlapping of spheres is allowed. Incurring huge performance cost.

The algorithm complexity can be measured.

- Time complexity $O(N^2)$: This can be seen as for every data point in the dataset, every data point is again visited to find its neighbour. Optimizations like using indexes and parallelisation are used to optimize the running time.
- Space complexity $O(N^2)$: For each data point all of its neighbours are stored. This overhead can be reduced by instead of storing the all the neighbours, whatever computation is needed must be done and stored and the cluster neighbours are thrown away. The only drawback is when doing a different computation, the neighbour will have to be found again.

Algorithm 2 Distance based clustering

```
1: procedure DISTANCE_BASED( $a, b$ )
2:   Input dataset: a list of data points, tol: tolerance, oTol: output tolerance
3:   Result subset of the data points that violate functional relation
4:   cluster  $\leftarrow$  Dictionary.empty
5:   for each point in dataset do ▷  $N$  iterations
6:     cluster[point]  $\leftarrow$  List.empty
7:     for each neighbour in dataset do ▷  $N$  iterations
8:       if  $\text{dist}(\text{point}, \text{neighbour}) < \text{tol}$  then
9:         cluster[point].add(neighbour)
10:  return cluster
```

2.6 Software

This section introduces the software and the framework which were chosen to implement the software. The software basic requirements were to be on of the mainstream programming language, open source, cross-platform and which has long-term community support. So the top options from which we had to choose from was Java, C/C++ and Python.

In [8], a very nice comparison is done between Java, Python and C. When it comes to portability, Java programs can compile to portable executable bytecode which can be run on any computer as long as it supports Java virtual machine. ANSI C, which can achieve the same portability require re-compilation of the source files. This means Java programs can be distributed as binary files that can run on any platform whereas C would require recompilation tools to recompile on the specific hardware. While Python code can enjoy this advantage and can run on any machine with a Python interpreter installed, the C-based Python extension modules, as well as the central Python interpreter itself, are only portable after (sometimes painful) recompilation of C source. With the ease of portability and cross-platform, we also know that Java programs are robust and can never have segmentation fault, and most of the errors are catchable runtime exceptions. On the other hand in C uncatchable, destructive errors all too frequently at runtime [8].

Garbage collection is supported by both Python and Java with the exception of C. Making it easier to code and build software. But since performance is another aspect, C and Java top Python because of the interpreted nature of the Python [8]. C is more performant than Java because of AOT (ahead of time compilation) in C than JIT (Just in time compilation) of Java. But writing and reading multi-threaded code is much easier in Java because of their nice set of portable API and well documentation. Python GIL (global interpreter lock) limits thread performance vastly [2].

Since we have to search through a lot of records, a lot of operation like grouping, summing, counting were required to be done on the dataset. Storing the dataset on memory using an array list, map or any data structure would be hard as well as error-prone with too much care needed for multi-threading. So, we decided to do the heavy dataset related operation from the database. Which is later be seen to use indexes to fasten up the search. With a little performance penalty, we would not have to worry about the memory limit anymore and disk space would be utilized by the database when not enough main memory is available.

A number of database models are out there namely relational model (e.g. MySQL), document model (e.g. MongoDB), graph model (e.g. Neo4j) and multi-model (e.g. ArangoDB). The dataset provided to the software can be with any number of columns to analyse. Dataset will be in the form of k columns separated by comma with each column containing “double” values where $k \in \mathbb{Z} \wedge k \geq 2$. So the storage of these datasets does not have a pre-defined schema. Grouping operation is one of the most important operations required. Also, we want to be able to represent the operation with a simple query language. ArangoDB stood out from all of the

databases, relational databases have really readable query language but its strict schema negates our requirement. In document model, MongoDB does support our dynamic model requirement but its query syntax is harder to read. Graph model is made for graph-based analysis whereas our analysis does not require graph. ArangoDB completes all of our requirement by supporting our dynamic model, having simple query language similar to MySQL named AQL and does support grouping operation providing graph operations as a bonus if ever required in the future.

With this, the conclusion was made to use Java as the programming language and ArangoDB as the database for hardcore data computing. Java has a large community support and ArangoDB fully supports Java client apps.

2.7 Applications

The reason for making a software like this verifies and gives the confidence about the dataset. If data do not fit any functional hypothesis in a space, much time could be saved by preventing the unneeded search. It also analysis the linear relation of the dataset. It assumes data is linear, does a linear regression on the various cluster and returns a multitude of values describing the data (e.g. number of clusters formed, number of outliers, mean of Pearson's coefficient R etc.).

It can be used to confirm the linear models already presented by [4] [3] [9]. The tool can be used to find the linear relationship in the dataset.

The software is not restricted to the use of only on the dataset which consists of performance events and power consumption. It is a general-purpose software which will for work for any kind of dataset in which user wants to know the existence of functional relation. The refuting of the claim of functional relation on the dataset is the objective of the software.

We also know that the dataset provided is usually experimentally measured values which are not accurate. Every measuring device has some margin of error. The software will be flexible in the sense that the equality comparison of values in the approaches will always be done keeping in the margin of error provided. It can also iterate certain of tolerance which will further be used to study the effect of tolerance on the dataset.

Chapter 3: Design and Optimisations

This section discusses the various details about software implemented in detail.

3.1 Design

The software is implemented in Java using ArangoDB as the backend database. The application has 3 layer structure as per Figure 3.1. The database layer which is the ArangoDB database itself. The business layer, this layer contains all the logic for talking to the backend ArangoDB, the clustering algorithms and other functions. The application has 2 frontends: Web application which is made using the Spring framework in Java and a simple command line application. Both the front end uses the same business layer hence providing the same functionality.

The way this project was strategised was to make command line first for easy testing of the core functionality like clustering, inserting the data, getting the data etc. Then the same core libraries are called via the web application. The web application was created for ease of use and multiple simultaneous analysis. Another reason for keeping the command line application was, some of the commands took too long to process. The connection between the server and the client used to timeout. Web sockets helped to keep the connection alive but once the connection is broken the results were lost in the process.

The whole project has 3 modules namely core, cmdapp, webapp. Gradle build tool was used for compiling, testing and running the application.




Figure 3.1: Structure of the app

3.1.1 ArangoDB

ArangoDB is a multi-model database. It supports three data models namely key/value, documents, graphs and a unified query language called AQL. AQL is very similar to SQL query

language. AQL is used for CRUD operations. It provides scalable and highly efficient queries. It uses JSON as a default storage schema.

The application uses this database for all the heavy duty data operations like grouping, summing, filtering neighbours etc. The reason a database was used instead of doing it directly in Java was for faster development, thread-safety and also the indexes provided by the database for faster queries.

3.1.2 Core

This business layer of the application contains all the core logic. It contains all the query, the driver to talk to ArangoDB, the logic for converting from and to CSV etc. The core module uses spring-context which is an application IoC (Inversion of Control) container. This container can be consumed by any application by adding the library and providing the database connection configuration. This saves the other application from creating each and every object. The IoC container creates the object for the application automatically at runtime. The strategy was to create a library which contains all main logic and which can be consumed by any kind of application. This can be seen as we have implemented in two applications command line and web application.

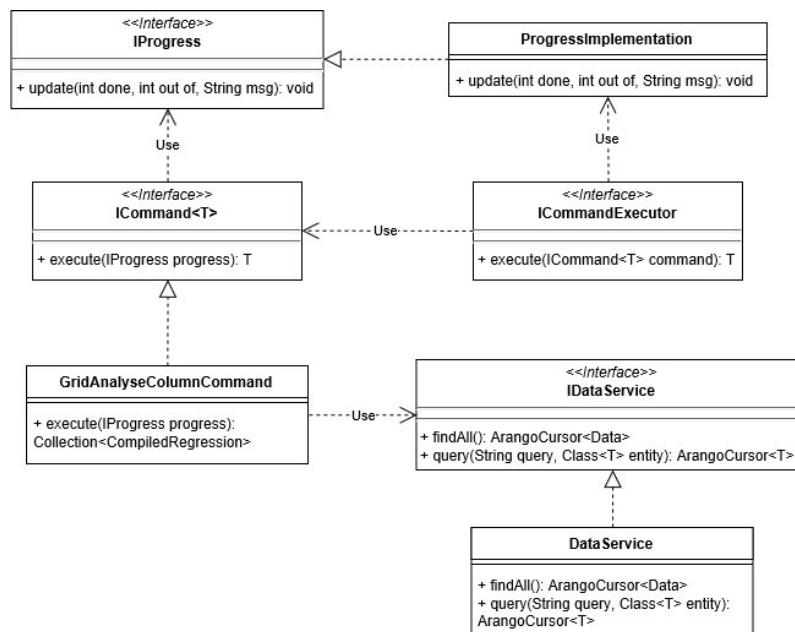


Figure 3.2: UML diagram of the command pattern in the core module

Apart from using IoC containers, the main design pattern used by this module of application is the Command Pattern [6]. In this pattern, an object is encapsulated with all the information needed to perform an action. The biggest feature of the command pattern is that it can be executed at any point in time. The command in our case is executed by a Command executor. This separation of the actual command and executor allows for greater changes in the future of the project. It provides a lot of flexibility and ease of adding and replacing commands in later stages of the development. One of the purposes of using command pattern was the certain routines in the application are quite slow e.g. distance-based clustering. The command pattern creates a unified way for getting the updates/progress of a task being executed.

The drawbacks of this pattern are it makes the source code larger quickly because one file can only represent one command. The number of files in the project increases very quickly. Command

names are usually big and so goes for the file names. The order of commands becomes important at later stages and can jeopardize the reliability of the system [6].

Figure 3.2, shows the implementation of command pattern. Only one Concrete command (GridAnalyseColumnCommand) is been shown. There are various commands implemented for data insertion, data deletion, distance-based analysing etc. It can be seen from the figure 3.2 that no implementation is defined for ICommandExecutor. This is because the modules consuming this core module library can implement according to their need. Similarly, for the IProgress, modules can implement their own version depending on their type of the application. The parameters of the command are supplied when creating the instance of the object. The command object constructor contains all the necessary arguments needed for execution at a later stage.

3.1.3 Webapp

The web application is built on Spring framework. It comes with IoC container which can easily pass in the database connection to the core library being consumed. In this module, the progress and command executor interfaces are implements to the needs of the web application. The progress interface implementation uses WebSocket for sending the update of how much work has been completed by the command. WebSocket is full-duplex communication channel over TCP which allows server and client to send messages to each other asynchronously.

The application maps the REST APIs to the Commands in the core module. The UI of the web application uses AngularJS which is a client-side front-end web framework. AngularJS framework allows to makes calls to the REST Endpoints and also act accordingly to the messages received by the server through WebSocket.

3.1.4 Command Line

The command line application also uses the core module to do a similar task. The mapping is from command line arguments to core module's command. In this module, as well a different implementation is done for the progress and command executor interfaces because this time the display and execution are in the command line. Since the core module uses IoC container. Spring-context was added and application context was created at the start of the application. The arguments were parsed via the JCommander library.

3.2 Optimizations

In this section, we discuss Optimizations that were done to speed up the computations.

3.2.1 Grid clustering algorithm

As per the pseudocode 1, it can be seen that the Time complexity is $O(N \log(N))$ which was due to the sorting performed on the coordinate of the box calculated for each data point. But as one can see from the pseudo code the two N iterations were performed. First, for calculating the coordinate of the k – cube they belong to and then for the grouping the data points into the

result. This result is then used for either finding violation of functional relation or for analysing the relation between the parameters.

The two extra N iterations were optimised by the ArangoDB by dividing the 3 iterations into 2 nodes. The first node sorts the data points by calculating the coordinates of the $k - cube$ on the go and attaching it temporarily to the dataset for reuse if the data set is accessed again by the sort. After sorting the data points with its coordinates are send to the second node. The second node then groups the data on the fly as the groups are accessed. This is because the all the groups with equal coordinate lies next to each other. The groups and the data coming from the second node is pipelined to our computations that are performed on the nodes.

It is impossible to parallelize this algorithm because of its aggregating nature, where one point cannot be in more than one cluster. But due to its $O(N \log(N))$ complexity, overall the algorithm is quite fast and linear.

3.2.2 Distance based clustering

In the distance based clustering 2, we calculated the time complexity to be $O(N^2)$. This was due to nature of the algorithm to find neighbouring data points for every data point in the dataset.

It is not possible to reduce the first outer iteration, as we have to go through each data point for analysing the whole dataset. The second one, it is not required to go through the entire data points as long as we can find all the points that are close to each other within some tolerance. This makes it possible for optimisations. One of the approaches is to save N^2 distances between each and every data point and then making it easier for all the subsequent queries. This is very practical and blazingly fast as only one time, a slow operation is required for saving all the distances between each and every point. Then finding neighbours is a trivial task of going through the entire edges and getting all those points where the distance is less than the specified tolerance. This is practical for smaller dataset but due to the square space complexity. But as the number of records gets larger it becomes inefficient to not only run the one-time insertion but also storing the same.

So, another way was chosen in which indexes were created to fasten the search to find the neighbours. We filter the data with conditions such that data points whose possibility of becoming neighbour is zero are filtered before the distance is actually and calculate and confirmed.

We use the fact that $dist(v, u) \geq 0$ always where v, u are two data points vectors in space.

Let the tol be the max distance (tolerance) allowed between two data points. Let u be the data point vector with dimension k whose neighbours are to be found. Let v be any vector in the dataset.

Then for any value of:

$$v[i] > u[i] + tol \vee v[i] < u[i] - tol$$

where $i \in [0, k - 1]$

$$\implies dist(u, v) > tol$$

The reason is if any one of the values in vector v is more than tol units away then because the distance is added from all the other values of the vector as well. The total distance will definitely be more than tol .

But this filtering method is still not good as we still have to go through all the dataset to check whether any of the value is greater than tol or not. ArangoDB provides APIs to create indexes

on the fields of documents for faster access to documents. It provides many different types of indexes such as Hash index, skip-list index, full-text index etc. Each index uses a special data structure to fasten different types of searches. As we are interested in range based queries. Skip list comes into play. Skiplist maintains a separate linked hierarchy of ordered values. They fasten up the equality lookups, range queries and sorting. The skip-list is added on to every parameter of the dataset and then when the filtering out condition is used before calculating the distance. The number of iteration will drop from N to the *Amortised* $- N$.

Another performance improvement that is done is by parallelising the outer loops. As ArangoDB supports multiple reads at the same time and the inner loop has no dependency on other iteration. The outer loop is parallelised which further decreases the running time by the number of threads the CPU can handle at a time.

3.3 Testing

All the functionality in core module has been well tested using the JUnit framework. Unit tests were created for classes which did not have external dependencies whereas Integration tests were created for classes that depended on other modules such as Commands, Data service etc. The reason unit test was not preferred because the unit test requires mocking of all of its dependencies. Integration test was preferred because it gave more confidence that the system works correctly as a whole. For components like analyse parameter etc, integration tests were created. Test data was entered directly into the database and commands were then called, asserting the results by the commands. Integration testing is not chosen for larger projects as these tests are very much slower than the unit test. But because our project involves a lot of database queries and is not at scale with industrial software, integration tests were performed.

Tests were created for almost every component and module. Rigorous testing was performed on core modules. As this module contains all the logic and interaction to the database. For parts such as clustering, regression, small records were created whose result we knew. These small records were entered into the database and clustering algorithms were called such that it gives the output known already. The testing is more of the black box testing form. A number of small hand-made datasets and their known output is used to test the components.

Chapter 4: Results and Evaluation

When the task is to find the existence of functional relation, Data points which violate the functional relation definition are returned. If there are no such data points no records are returned. But when we move on to next step of analysing the functional relation a number of variables describing the clusters formed for analysing a particular parameter variable are returned. Following is the list of those variables.

- *Parameter nos*: index of the parameter that the analyses are run for.
- *Tolerance*: the tolerance which was given as input to be used for the clustering algorithm.
- *Slope of the regression line*: mean and standard deviation of the slope from all the clusters formed.
- *Y-intercept of the regression line*: mean and standard deviation of the y-intercept from all the clusters.
- *Weighted slope and y-intercept*: Since finding a regression line requires only minimum 2 points. A number of small clusters can hide the values of the dense clusters. Weighted slope and y-intercept multiply the values with the number of points in the data cluster to give more weightage to the dense clusters.
- *Pearson correlation coefficient*: It is the measure of the linear correlation between two variables X and Y. It ranges from -1 to 1 represent total negative correlation to total positive correlation.
- R^2 : it is coefficient of determination explaining how good is the regression line fit.
- *Number of outliers*: these are the numbers of clusters formed with only 1 data point.
- *Number of cluster*: this represents the number of clusters formed with greater than 1 data member.
- *Number of data points in each cluster*: average and standard deviation of the number of data points in each cluster is returned
- Number of points that failed to provide a linear regression line: there are times when the linear regression cannot be calculated because the data points in the cluster have equal non clustered parameter variable values.

4.1 Results

The software was run on computer-generated linear dataset as well as the real-life experimental dataset. And for computer-generated records. We found that the value of the *constants* used in the generation of the records was equivalent to the mean of the *slope* computed from all the clusters formed with the accuracy of 10%. The accuracy increased if the *constants* were larger and decreased if the *constants* were smaller. This was because it is easy to see the change when the parameter's contribution to the output is large. The *standard deviation of the slope*

was low when the function was linear. If any of the parameters contributions was non-linear e.g. polynomial, the *standard deviation of the slope* increased significantly. *Weighted values* always had higher accuracy than the non-weighted values.

The software requires an argument namely tolerance which is used as a threshold for clustering. Below Figure 4.1 shows the mean of Pearson coefficient R when the tolerance is increased form 0% to 12% for the experimental small dataset provided for the parameter number 4. Pearson coefficient is the measure of linear correlation. It tells how strong is the correlation between two variables. It can be seen in the graph for very low tolerance the Pearson coefficient is unreliable. This is because not many clusters are formed properly yet. As soon as the tolerance is increased. The correlation value increases substantially and reaches a maximum point for certain tolerance. After that, a further increase in tolerance degrades the correlation. This is because non-neighbour points also get clustered due to the tolerance factor.

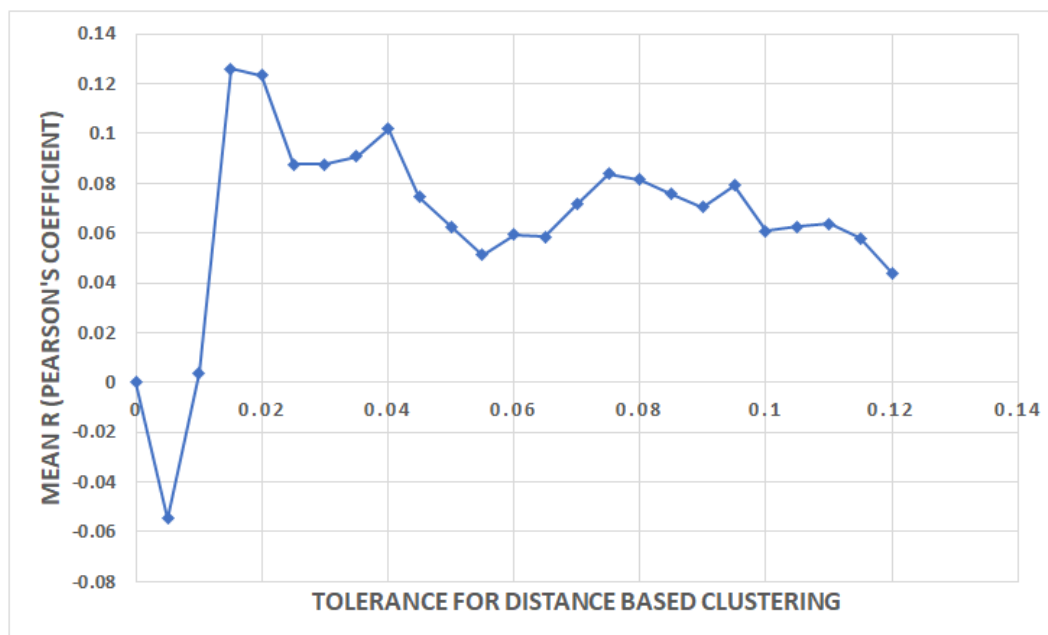


Figure 4.1: Pearson Coefficient vs Tolerance

4.2 Challenges

One of the challenges that faced while doing the analysis of the dataset was when the points are too far away from each other, this was because the units of one of the variable were really large. Since the parameter variable units were not comparable. The clustering algorithm was failing. Hence, to make it comparable the data was normalized. This normalization helped in making the parameters comparable. The tolerance was adjusted accordingly.

Another challenge faced was that because of normalization the real value of the slope got hidden. The normalization used was feature scaling the scales were changed and hence the slope and other values of the dataset did not correspond to what they should. To solve this problem, the dataset was then clustered using normalized value but the regression line was created for the raw values of the dataset. This help solves the clustering as well as the regression problem.

Chapter 5: Conclusions and Future Work

5.1 Conclusion

In this project, a successful software was created which enabled the ability to understand the relation between input parameter variable and the output variable. The software was run on experimental data set where the input parameters and the output variable corresponds to the performance events and the energy consumption. The software enabled us to cluster the data with various tolerance and analyse the effect of tolerance. This also helps in understanding and finding the optimum tolerance in the dataset where the functional relation is linear the most. It can be used for validating the linear models already in place and understand the relationship between variables on a much higher level.

The project has certain limitation such as the presence of a non-influential performance event or parameter could lead inaccurate result. The software works well with noise as the first step of the existence of functional relation could stop the noise from the analysis. But the addition of noise variable would corrupt the results significantly. Another limitation of the project is that it can only verify linear models and can only analyse linear relation.

5.2 Future work

More work is needed in the clustering algorithm because the Euclidean distance is not always the right choice for finding the difference between two points. As per [1], the Euclidean distance measure does not work accurately because of the effect of the curse of high dimensionality.

The regression analysis right now is only linear. We would want to be able to analyse and validate other models such as log, polynomial etc. Each cluster must go through different regression analysis and pick the model which best fits the dataset.

Bibliography

- [1] Charu C Aggarwal, Alexander Hinneburg, and Daniel A Keim. On the surprising behavior of distance metrics in high dimensional space. In *International conference on database theory*, pages 420–434. Springer, 2001.
- [2] David Beazley. Understanding the python gil. In *PyCON Python Conference. Atlanta, Georgia*, 2010.
- [3] Frank Bellosa. The benefits of event: driven energy accounting in power-sensitive systems. In *Proceedings of the 9th workshop on ACM SIGOPS European workshop: beyond the PC: new challenges for the operating system*, pages 37–42. ACM, 2000.
- [4] W Lloyd Bircher and Lizy K John. Complete system power estimation: A trickle-down approach based on performance events. In *Performance Analysis of Systems & Software, 2007. ISPASS 2007. IEEE International Symposium on*, pages 158–168. IEEE, 2007.
- [5] William Lloyd Bircher and Lizy K John. Complete system power estimation using processor performance events. *IEEE Transactions on Computers*, 61(4):563–577, 2012.
- [6] Benoit Dupire and Eduardo B Fernandez. The command dispatcher pattern. In *8th Conference on Pattern Languages of Programs (PLOP 01)*, 2001.
- [7] C Gilberto and M Margaret. Power prediction for intel xscale processors using performance monitoring unit events power prediction for intel xscale processors using performance monitoring unit events. In *ISLPED*, volume 5, pages 8–10, 2005.
- [8] Jim Hugunin. Python and java: The best of both worlds. In *Proceedings of the 6th international Python conference*, volume 9, pages 2–18, 1997.
- [9] Kenneth Obrien, Ilia Pietri, Ravi Reddy, Alexey Lastovetsky, and Rizos Sakellariou. A survey of power and energy predictive models in hpc systems and applications. *ACM Computing Surveys (CSUR)*, 50(3):37, 2017.
- [10] Mario Pickavet, Willem Vereecken, Sofie Demeyer, Pieter Audenaert, Brecht Vermeulen, Chris Develder, Didier Colle, Bart Dhoedt, and Piet Demeester. Worldwide energy needs for ict: The rise of power-aware networking. In *Advanced Networks and Telecommunication Systems, 2008. ANTS'08. 2nd International Symposium on*, pages 1–3. IEEE, 2008.
- [11] Huigui Rong, Haomin Zhang, Sheng Xiao, Canbing Li, and Chunhua Hu. Optimizing energy consumption for data centers. *Renewable and Sustainable Energy Reviews*, 58:674–691, 2016.
- [12] Rui Xu and Donald Wunsch. Survey of clustering algorithms. *IEEE Transactions on neural networks*, 16(3):645–678, 2005.
- [13] Shu Yang, Zhongzhi Luan, Binyang Li, Ge Zhang, Tianming Huang, and Depei Qian. Performance events based full system estimation on application power consumption. In *High Performance Computing and Communications; IEEE 14th International Conference on Smart City; IEEE 2nd International Conference on Data Science and Systems (HPCC/SmartCity/DSS), 2016 IEEE 18th International Conference on*, pages 749–756. IEEE, 2016.
- [14] Robert Zembowicz and Jan M Zytkow. Testing the existence of functional relationship in data. In *Proceedings of AAAI Workshop on Knowledge Discovery in Databases*, volume 93, pages 102–111, 1993.

Appendix A:

Project repository link: <https://github.com/Sukrat/FYP2017-function-relationship-finder>